

## Module 3 Lab Exercise: Machine Learning Workflow and Types of Learning

### Learning Objectives

By the end of this lab, you will be able to:

- Distinguish between supervised, unsupervised, and reinforcement learning
- Understand the complete machine learning workflow
- Build and evaluate your first classification model
- Work with different types of data (numerical, categorical, text, images)
- Apply the end-to-end ML process: data → model → evaluation → insights

### Prerequisites

- Completed Module 2 (familiar with Python libraries and Jupyter/Colab)
- Understanding of basic data operations and visualization
- Access to your GitHub repository for saving work

### Part 1: Understanding Types of Machine Learning

Machine learning can be categorized into three main types. Let's explore each with practical examples.

#### 1. Supervised Learning

**Definition:** Learning from labeled examples to make predictions on new, unseen data.

**Examples:**

- **Classification:** Predicting categories (spam/not spam, disease/healthy)
- **Regression:** Predicting continuous values (house prices, temperature)

**Key Characteristic:** We have both input features (X) and correct answers (y) during training.

#### 2. Unsupervised Learning

**Definition:** Finding hidden patterns in data without labeled examples.

**Examples:**

- **Clustering:** Grouping similar customers for marketing
- **Dimensionality Reduction:** Simplifying complex data while keeping important information

**Key Characteristic:** We only have input features (X), no correct answers during training.

#### 3. Reinforcement Learning

**Definition:** Learning through trial and error by receiving rewards or penalties.

**Examples:**

- Game playing (chess, Go)
- Autonomous vehicles
- Recommendation systems that learn from user feedback

**Key Characteristic:** Agent learns by interacting with an environment and receiving feedback.

For this course, we'll focus primarily on supervised learning, with some unsupervised learning in later modules.

## Part 2: Setting Up Our Machine Learning Environment

Let's start by importing our libraries and loading a dataset that will help us understand the ML workflow.

```
# Import essential libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_wine, make_classification
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

```

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')

# Set style for better-looking plots
plt.style.use('default')
sns.set_palette("husl")

print("✅ All libraries imported successfully!")
print("🚀 Ready to start our machine learning journey!")

✅ All libraries imported successfully!
🚀 Ready to start our machine learning journey!

```

## Part 3: Loading and Exploring Our Dataset

We'll use the Wine dataset - a classic dataset for classification. It contains chemical analysis of wines from three different cultivars (types) grown in Italy.

```

# Load the Wine dataset
wine_data = load_wine()

# Convert to DataFrame for easier handling
df = pd.DataFrame(wine_data.data, columns=wine_data.feature_names)
df['wine_class'] = wine_data.target
df['wine_class_name'] = [wine_data.target_names[i] for i in wine_data.target]

print("Dataset Information:")
print(f"Shape: {df.shape}")
print(f"Features: {len(wine_data.feature_names)}")
print(f"Classes: {wine_data.target_names}")
print("\nFirst 5 rows:")
print(df.head())

Dataset Information:
Shape: (178, 15)
Features: 13
Classes: ['class_0' 'class_1' 'class_2']

First 5 rows:
   alcohol  malic_acid    ash  alcalinity_of_ash  magnesium  total_phenols \
0      14.23       1.71  2.43                  15.6        127.0          2.80
1      13.20       1.78  2.14                  11.2        100.0          2.65
2      13.16       2.36  2.67                  18.6        101.0          2.80
3      14.37       1.95  2.50                  16.8        113.0          3.85
4      13.24       2.59  2.87                  21.0        118.0          2.80

   flavanoids  nonflavanoid_phenols  proanthocyanins  color_intensity    hue \
0         3.06                 0.28            2.29           5.64  1.04
1         2.76                 0.26            1.28           4.38  1.05
2         3.24                 0.30            2.81           5.68  1.03
3         3.49                 0.24            2.18           7.80  0.86
4         2.69                 0.39            1.82           4.32  1.04

   od280/od315_of_diluted_wines  proline  wine_class wine_class_name
0                  3.92    1065.0       0     class_0
1                  3.40    1050.0       0     class_0
2                  3.17    1185.0       0     class_0
3                  3.45    1480.0       0     class_0
4                  2.93     735.0       0     class_0

```

```

# Explore the dataset structure
print("Dataset Overview:")
print("=" * 50)
print(f"Total samples: {len(df)}")
print(f"Features (input variables): {len(df.columns) - 2}") # -2 for target columns
print(f"Target classes: {df['wine_class_name'].unique()}")
print("\nClass distribution:")
print(df['wine_class_name'].value_counts())

# Check for missing values
print("\nMissing values: {df.isnull().sum().sum()}")
print("✅ No missing values - this is a clean dataset!")

```

```

Dataset Overview:
=====
Total samples: 178
Features (input variables): 13
Target classes: [np.str_('class_0') np.str_('class_1') np.str_('class_2')]

Class distribution:

```

```
wine_class_name
class_1    71
class_0    59
class_2    48
Name: count, dtype: int64

Missing values: 0
 No missing values - this is a clean dataset!
```

## Part 4: Exploratory Data Analysis (EDA)

Before building models, we need to understand our data. This is a crucial step in the ML workflow.



```
# Visualize class distribution
plt.figure(figsize=(12, 4))

# Subplot 1: Class distribution
plt.subplot(1, 2, 1)
class_counts = df['wine_class_name'].value_counts()
plt.bar(class_counts.index, class_counts.values, color=['red', 'green', 'blue'])
plt.title('Distribution of Wine Classes')
plt.xlabel('Wine Class')
plt.ylabel('Number of Samples')
plt.xticks(rotation=45)

# Subplot 2: Feature correlation heatmap (first 6 features for clarity)
plt.subplot(1, 2, 2)
correlation_matrix = df.iloc[:, :6].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0, fmt='.2f')
plt.title('Feature Correlations (First 6 Features)')

plt.tight_layout()
plt.show()

print("📊 EDA helps us understand:")
print("- Class balance (are all classes equally represented?)")
print("- Feature relationships (which features are correlated?)")
print("- Data quality (any outliers or issues?)")
```



📊 EDA helps us understand:  
- Class balance (are all classes equally represented?)  
- Feature relationships (which features are correlated?)  
- Data quality (any outliers or issues?)

## Part 5: The Complete Machine Learning Workflow

Now let's implement the standard ML workflow step by step:

The 6-Step ML Workflow:

1. **Data Preparation:** Clean and prepare the data
2. **Feature Selection:** Choose relevant input variables
3. **Data Splitting:** Separate training and testing data
4. **Model Training:** Teach the algorithm using training data
5. **Model Evaluation:** Test performance on unseen data
6. **Model Interpretation:** Understand what the model learned

Let's implement each step!

```
# Step 1: Data Preparation
print("Step 1: Data Preparation")
print("=" * 30)

# Select features (X) and target (y)
# For simplicity, let's use the first 4 features
feature_names = ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash']
X = df[feature_names]
y = df['wine_class']

print(f"Selected features: {feature_names}")
print(f"Feature matrix shape: {X.shape}")
print(f"Target vector shape: {y.shape}")

# Display first few rows
print("\nFirst 5 samples:")
print(X.head())
```

Step 1: Data Preparation  
=====

```
Selected features: ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash']
Feature matrix shape: (178, 4)
Target vector shape: (178,)

First 5 samples:
   alcohol  malic_acid    ash  alcalinity_of_ash
0      14.23        1.71  2.43          15.6
1      13.20        1.78  2.14          11.2
2      13.16        2.36  2.67          18.6
3      14.37        1.95  2.50          16.8
4      13.24        2.59  2.87          21.0
```

```
# Step 2: Data Splitting
print("Step 2: Data Splitting")
print("=" * 30)

# Split data into training (80%) and testing (20%) sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,      # 20% for testing
    random_state=42,    # For reproducible results
    stratify=y          # Maintain class proportions
)

print(f"Training set: {X_train.shape[0]} samples")
print(f"Testing set: {X_test.shape[0]} samples")
print(f"Training classes: {np.bincount(y_train)}")
print(f"Testing classes: {np.bincount(y_test)}")

print("\n⌚ Why split data?")
print("- Training set: Teach the model")
print("- Testing set: Evaluate performance on unseen data")
print("- This prevents overfitting (memorizing vs. learning)")
```

Step 2: Data Splitting  
=====

```
Training set: 142 samples
Testing set: 36 samples
Training classes: [47 57 38]
Testing classes: [12 14 10]

⌚ Why split data?
- Training set: Teach the model
- Testing set: Evaluate performance on unseen data
- This prevents overfitting (memorizing vs. learning)
```

```
# Step 3: Model Training
print("Step 3: Model Training")
print("=" * 30)

# Create and train two different models
models = {
    'Logistic Regression': LogisticRegression(random_state=42),
    'Decision Tree': DecisionTreeClassifier(random_state=42, max_depth=3)
}

trained_models = {}

for name, model in models.items():
    print(f"\nTraining {name}...")
```

```

# Train the model
model.fit(X_train, y_train)
trained_models[name] = model

print(f"✅ {name} training completed!")

print("\n💡 What happened during training?")
print("- Models learned patterns from training data")
print("- They found relationships between features and wine classes")
print("- Now they can make predictions on new data!")

```

Step 3: Model Training  
=====

Training Logistic Regression...  
✅ Logistic Regression training completed!

Training Decision Tree...  
✅ Decision Tree training completed!

💡 What happened during training?  
- Models learned patterns from training data  
- They found relationships between features and wine classes  
- Now they can make predictions on new data!

```

# Step 4: Model Evaluation
print("Step 4: Model Evaluation")
print("=" * 30)

results = {}

for name, model in trained_models.items():
    # Make predictions
    y_pred = model.predict(X_test)

    # Calculate accuracy
    accuracy = accuracy_score(y_test, y_pred)
    results[name] = accuracy

    print(f"\n{name} Results:")
    print(f"Accuracy: {accuracy:.3f} ({accuracy*100:.1f}%)")

    # Detailed classification report
    print("\nDetailed Performance:")
    print(classification_report(y_test, y_pred, target_names=wine_data.target_names))

# Compare models
print("\n📊 Model Comparison:")
for name, accuracy in results.items():
    print(f"{name}: {accuracy:.3f}")

best_model = max(results, key=results.get)
print(f"\n🏆 Best performing model: {best_model}")

```

Step 4: Model Evaluation  
=====

Logistic Regression Results:  
Accuracy: 0.889 (88.9%)

Detailed Performance:

	precision	recall	f1-score	support
class_0	1.00	1.00	1.00	12
class_1	0.81	0.93	0.87	14
class_2	0.88	0.70	0.78	10
accuracy			0.89	36
macro avg	0.90	0.88	0.88	36
weighted avg	0.89	0.89	0.89	36

Decision Tree Results:  
Accuracy: 0.833 (83.3%)

Detailed Performance:

	precision	recall	f1-score	support
class_0	0.86	1.00	0.92	12
class_1	0.91	0.71	0.80	14
class_2	0.73	0.80	0.76	10
accuracy			0.83	36
macro avg	0.83	0.84	0.83	36
weighted avg	0.84	0.83	0.83	36

Model Comparison:  
 Logistic Regression: 0.889  
 Decision Tree: 0.833

Best performing model: Logistic Regression


```
# Step 5: Model Interpretation
print("Step 5: Model Interpretation")
print("=" * 30)

# Visualize confusion matrix for the best model
best_model_obj = trained_models[best_model]
y_pred_best = best_model_obj.predict(X_test)

plt.figure(figsize=(8, 6))
cm = confusion_matrix(y_test, y_pred_best)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=wine_data.target_names,
            yticklabels=wine_data.target_names)
plt.title(f'Confusion Matrix - {best_model}')
plt.xlabel('Predicted Class')
plt.ylabel('True Class')
plt.show()

print(f"\n🔍 Interpreting the Confusion Matrix:")
print("- Diagonal values: Correct predictions")
print("- Off-diagonal values: Misclassifications")
print("- Perfect model would have all values on diagonal")
```

Step 5: Model Interpretation



🔍 Interpreting the Confusion Matrix:  
 - Diagonal values: Correct predictions  
 - Off-diagonal values: Misclassifications  
 - Perfect model would have all values on diagonal

## Part 6: Understanding Different Data Types in ML

Machine learning works with various types of data. Let's explore the main categories:

```
# Understanding Different Data Types in ML
print("Understanding Data Types in Machine Learning")
print("=" * 45)

# Create examples of different data types
data_examples = {
    'Numerical (Continuous)': [23.5, 45.2, 67.8, 12.1, 89.3],
    'Numerical (Discrete)': [1, 5, 3, 8, 2],
    'Categorical (Nominal)': ['Red', 'Blue', 'Green', 'Red', 'Blue'],
```

```
'Categorical (Ordinal)': ['Low', 'Medium', 'High', 'Medium', 'Low'],
'Text': ['Hello world', 'Machine learning', 'Data science', 'Python programming', 'AI revolution'],
'Boolean': [True, False, True, True, False]
}

for data_type, examples in data_examples.items():
    print(f"\n{data_type}:")
    print(f"  Examples: {examples}")
    print(f"  Use case: ", end="")

    if 'Continuous' in data_type:
        print("Regression problems (predicting prices, temperatures)")
    elif 'Discrete' in data_type:
        print("Counting problems (number of items, ratings)")
    elif 'Nominal' in data_type:
        print("Classification without order (colors, categories)")
    elif 'Ordinal' in data_type:
        print("Classification with order (ratings, sizes)")
    elif 'Text' in data_type:
        print("Natural language processing (sentiment analysis, translation)")
    elif 'Boolean' in data_type:
        print("Binary classification (yes/no, spam/not spam)")

print("\n💡 Key Insight: Different data types require different preprocessing and algorithms!")
```

### Understanding Data Types in Machine Learning

---

#### Numerical (Continuous):

Examples: [23.5, 45.2, 67.8, 12.1, 89.3]  
Use case: Regression problems (predicting prices, temperatures)

#### Numerical (Discrete):

Examples: [1, 5, 3, 8, 2]  
Use case: Counting problems (number of items, ratings)

#### Categorical (Nominal):

Examples: ['Red', 'Blue', 'Green', 'Red', 'Blue']  
Use case: Classification without order (colors, categories)

#### Categorical (Ordinal):

Examples: ['Low', 'Medium', 'High', 'Medium', 'Low']  
Use case: Classification with order (ratings, sizes)

#### Text:

Examples: ['Hello world', 'Machine learning', 'Data science', 'Python programming', 'AI revolution']  
Use case: Natural language processing (sentiment analysis, translation)

#### Boolean:

Examples: [True, False, True, True, False]  
Use case: Binary classification (yes/no, spam/not spam)

💡 Key Insight: Different data types require different preprocessing and algorithms!

## Part 7: Hands-On Practice - Build Your Own Model

Now it's your turn! Complete the following tasks to reinforce your learning.

```
# Task 1: Try different features
print("Task 1: Experiment with Different Features")
print("=" * 40)

# Your task: Select 3 different features and build a model
# Available features:
print("Available features:")
for i, feature in enumerate(wine_data.feature_names):
    print(f"{i+1}: {feature}")

# TODO: Replace these with your chosen features
your_features = ['alcohol', 'color_intensity', 'proline'] # Modify this list

# Build model with your features
X_your = df[your_features]
X_train_your, X_test_your, y_train_your, y_test_your = train_test_split(
    X_your, y, test_size=0.2, random_state=42, stratify=y
)

# Train a logistic regression model
your_model = LogisticRegression(random_state=42)
your_model.fit(X_train_your, y_train_your)

# Evaluate
# pred_your = your_model.predict(X_test_your)
```

```

your_accuracy = accuracy_score(y_test_your, y_pred_your)

print(f"\nYour model features: {your_features}")
print(f"Your model accuracy: {your_accuracy:.3f} ({your_accuracy*100:.1f}%)")

# Compare with original model
print(f"Original model accuracy: {results['Logistic Regression']:.3f}")
if your_accuracy > results['Logistic Regression']:
    print("🎉 Great job! Your feature selection improved the model!")
else:
    print("💡 Try different features to see if you can improve performance!")

```

#### Task 1: Experiment with Different Features

---

Available features:

1. alcohol
2. malic\_acid
3. ash
4. alcalinity\_of\_ash
5. magnesium
6. total\_phenols
7. flavanoids
8. nonflavanoid\_phenols
9. proanthocyanins
10. color\_intensity
11. hue
12. od280/od315\_of\_diluted\_wines
13. proline

Your model features: ['alcohol', 'color\_intensity', 'proline']

Your model accuracy: 0.833 (83.3%)

Original model accuracy: 0.889

💡 Try different features to see if you can improve performance!

## ▼ Part 8: Assessment - Understanding ML Concepts

Answer the following questions to demonstrate your understanding:

```

# Assessment Task 1: Identify the ML type
print("Assessment Task 1: Identify Machine Learning Types")
print("=" * 50)

# For each scenario, identify if it's Supervised, Unsupervised, or Reinforcement Learning

scenarios = [
    "Predicting house prices based on size, location, and age",
    "Grouping customers by purchasing behavior without knowing groups beforehand",
    "Teaching a robot to play chess by playing many games",
    "Classifying emails as spam or not spam using labeled examples",
    "Finding hidden topics in news articles without predefined categories"
]

# Your answers (replace 'TYPE' with Supervised, Unsupervised, or Reinforcement)
your_answers = [
    "Supervised",      # Scenario 1
    "Unsupervised",   # Scenario 2
    "Reinforcement",  # Scenario 3
    "Supervised",      # Scenario 4
    "Unsupervised"     # Scenario 5
]

# Check answers
correct_answers = ["Supervised", "Unsupervised", "Reinforcement", "Supervised", "Unsupervised"]

print("Scenario Analysis:")
score = 0
for i, (scenario, your_answer, correct) in enumerate(zip(scenarios, your_answers, correct_answers)):
    is_correct = your_answer == correct
    score += is_correct
    status = "✅" if is_correct else "❌"
    print(f"{status} {i+1}. {scenario}")
    print(f"  Your answer: {your_answer} | Correct: {correct}")
    print()

print(f"Score: {score}/{len(scenarios)} ({score/len(scenarios)*100:.0f}%)")

```

Assessment Task 1: Identify Machine Learning Types

---

Scenario Analysis:

- 1. Predicting house prices based on size, location, and age

Your answer: Supervised | Correct: Supervised

- 2. Grouping customers by purchasing behavior without knowing groups beforehand  
Your answer: Unsupervised | Correct: Unsupervised
- 3. Teaching a robot to play chess by playing many games  
Your answer: Reinforcement | Correct: Reinforcement
- 4. Classifying emails as spam or not spam using labeled examples  
Your answer: Supervised | Correct: Supervised
- 5. Finding hidden topics in news articles without predefined categories  
Your answer: Unsupervised | Correct: Unsupervised

Score: 5/5 (100%)

## Part 9: Real-World Applications and Case Studies

Let's explore how the concepts we've learned apply to real-world scenarios.

### Case Study 1: Recommendation Systems (Netflix, Amazon)

**Problem:** Suggest movies/products users might like **ML Type:** Hybrid (Supervised + Unsupervised + Reinforcement) **Data:** User ratings, viewing history, product features **Workflow:** Collect data → Build user profiles → Train models → Make recommendations → Learn from feedback

### Case Study 2: Fraud Detection (Banks, Credit Cards)

**Problem:** Identify fraudulent transactions **ML Type:** Supervised Learning (Classification) **Data:** Transaction amounts, locations, times, merchant types **Workflow:** Historical fraud data → Feature engineering → Train classifier → Real-time scoring → Continuous monitoring

### Case Study 3: Medical Diagnosis (Healthcare)

**Problem:** Assist doctors in diagnosing diseases **ML Type:** Supervised Learning (Classification) **Data:** Medical images, patient symptoms, lab results **Workflow:** Labeled medical data → Image processing → Train deep learning models → Clinical validation → Deployment with human oversight

### Your Turn: Think of Applications

Consider these industries and think about how ML could be applied:

- **Transportation:** Autonomous vehicles, route optimization
- **Agriculture:** Crop monitoring, yield prediction
- **Education:** Personalized learning, automated grading
- **Entertainment:** Content creation, game AI

## Part 10: Complete ML Workflow Summary

Let's summarize the complete machine learning workflow we've learned:

### The Machine Learning Lifecycle

1. Problem Definition  
↓
2. Data Collection & Exploration  
↓
3. Data Preprocessing & Feature Engineering  
↓
4. Model Selection & Training  
↓
5. Model Evaluation & Validation  
↓
6. Model Deployment & Monitoring  
↓
7. Continuous Improvement

### Checklist for Every ML Project:

#### Data Phase:

- Understand the problem and define success metrics
- Collect and explore the dataset
- Check for missing values, outliers, and data quality issues

- Visualize data to understand patterns and relationships

#### **Modeling Phase:**

- Split data into training and testing sets
- Select appropriate algorithms for the problem type
- Train multiple models and compare performance
- Evaluate using appropriate metrics (accuracy, precision, recall, etc.)

#### **Deployment Phase:**

- Validate model performance on new data
- Document the model and its limitations
- Deploy responsibly with monitoring systems
- Plan for model updates and maintenance

#### Key Takeaways:

1. **Start Simple:** Begin with basic models before trying complex ones
2. **Understand Your Data:** EDA is crucial for success
3. **Validate Properly:** Always test on unseen data
4. **Iterate:** ML is an iterative process of improvement
5. **Document Everything:** Keep track of experiments and results

## Your Reflection and Analysis

**Instructions:** Complete the reflection below by editing this markdown cell.

### My Understanding of Machine Learning Types

**Supervised Learning:** Its uses labeled data to train models for prediction.

**Unsupervised Learning:** Its uses unlabeled data to find hidden pattern and structure.

**Reinforcement Learning:** RL is a branch of machine learning in which models learn from errors and by collecting experiences, taking actions and observing what happens.

### My Analysis of the Wine Classification Project

**Best performing model:** Logistic Regression

**Why do you think this model performed better?**: Because the dataset is clean and small. Also, feature are numeric, it's doesn't have complex structure.

**What would you try next to improve performance?**: Advanced feature engineering, use non-linear transformation, tune decision threshold, dimensionality reduction and different representation of categorical data.

### Real-World Application Ideas

**Industry of Interest:** Predicting Traffic Congestion

**ML Problem:** Forecast traffic level

**Type of ML:** Supervised learning

**Data Needed:** Historical traffic volume, weather, road sensors and time.

### Key Learnings

**Most important concept learned:** Evaluated Model Performance

**Most challenging part:** Data splitting

**Questions for further exploration:** No

## Lab Summary and Next Steps

#### What You've Accomplished:

- ✓ **Understood ML Types:** Supervised, Unsupervised, and Reinforcement Learning
- ✓ **Mastered ML Workflow:** Data → Model → Evaluation → Insights
- ✓ **Built Classification Models:** Logistic Regression and Decision Trees
- ✓ **Evaluated Model Performance:** Accuracy, Confusion Matrix, Classification Report

- Worked with Real Data:** Wine dataset analysis and modeling
- Applied Best Practices:** Data splitting, model comparison, interpretation

### 🚀 Preparation for Module 4:

In the next lab, you'll dive deeper into:

- **Exploratory Data Analysis (EDA):** Advanced visualization techniques
- **Data Quality Assessment:** Handling missing values, outliers, and duplicates
- **Statistical Analysis:** Understanding distributions and relationships
- **Data Storytelling:** Communicating insights effectively

### 📝 Action Items:

1. **Upload this notebook** to your GitHub repository
2. **Experiment** with different features in the wine dataset
3. **Try other datasets** from sklearn.datasets (digits, breast\_cancer, boston)
4. **Practice** the 6-step ML workflow on a new problem
5. **Document** your experiments and findings

### 🔗 Additional Resources:

- [Scikit-learn User Guide](#)
- [Machine Learning Mastery](#)
- [Kaggle Learn](#) - Free micro-courses
- [Google's Machine Learning Crash Course](#)

### 🧠 Reflection Questions:

1. Which type of machine learning (supervised/unsupervised/reinforcement) interests you most and why?
2. What was the most challenging part of the ML workflow for you?
3. How might you apply these concepts to a problem in your field of interest?
4. What questions do you have about machine learning that you'd like to explore further?

**Congratulations on completing Module 3! You've taken a significant step in your machine learning journey.** 🎉

---

*Remember: Machine learning is a skill that improves with practice. Keep experimenting, stay curious, and don't be afraid to make mistakes - they're part of the learning process!*